



```

1 # Create a linear regression model in PyTorch
2 class LinearRegressionModel(nn.Module):
3     def __init__(self):
4         super().__init__()
5
6         # Initialize model parameters
7         self.weights = nn.Parameter(torch.randn(1),
8                                       requires_grad=True,
9                                       dtype=torch.float)
10
11         self.bias = nn.Parameter(torch.randn(1),
12                                  requires_grad=True,
13                                  dtype=torch.float)
14
15 # forward() defines the computation in the model
16 def forward(self, x: torch.Tensor) -> torch.Tensor:
17     return self.weights * x + self.bias

```

Subclass `nn.Module` (this contains all the building blocks for neural networks)

Initialize **model parameters** to be used in various computations (these could be different layers from `torch.nn`, single parameters, hard-coded values or functions)

`requires_grad=True` means PyTorch will track the gradients of this specific parameter for use with `torch.autograd` and gradient descent (for many `torch.nn` modules, `requires_grad=True` is set by default)

Any subclass of `nn.Module` needs to override `forward()` (this defines the forward computation of the model)

```

1 # Create a linear regression model in PyTorch
2 class LinearRegressionModel(nn.Module):
3     def __init__(self):
4         super().__init__()
5
6         # Initialize model parameters
7         self.weights = nn.Parameter(torch.randn(1),
8                                       requires_grad=True,
9                                       dtype=torch.float)
10
11         self.bias = nn.Parameter(torch.randn(1),
12                                  requires_grad=True,
13                                  dtype=torch.float)
14
15 # forward() defines the computation in the model
16 def forward(self, x: torch.Tensor) -> torch.Tensor:
17     return self.weights * x + self.bias

```

Linear regression model with `nn.Parameter`

```

1 # Create a linear regression model in PyTorch with nn.Linear
2 class LinearRegressionModel(nn.Module):
3     def __init__(self):
4         super().__init__()
5         # Use nn.Linear() for creating the model parameters
6         self.linear_layer = nn.Linear(in_features=1,
7                                       out_features=1)
8
9 # forward() defines the computation in the model
10 def forward(self, x: torch.Tensor) -> torch.Tensor:
11     return self.linear_layer(x)

```

Linear regression model with `nn.Linear`

# PyTorch training loop

```

1 # Pass the data through the model for a number of epochs (e.g. 100)
2 for epoch in range(epochs):
3     # Put model in training mode (this is the default state of a model)
4     model.train()
5
6     # 1. Forward pass on train data using the forward() method inside
7     y_pred = model(X_train)
8
9     # 2. Calculate the loss (how different are the model's predictions to the true values)
10    loss = loss_fn(y_pred, y_true)
11
12    # 3. Zero the gradients of the optimizer (they accumulate by default)
13    optimizer.zero_grad()
14
15    # 4. Perform backpropagation on the loss
16    loss.backward()
17
18    # 5. Progress/step the optimizer (gradient descent)
19    optimizer.step()
20

```

Pass the data through the model for a number of **epochs** (e.g. 100 for 100 passes of the data)

Pass the data through the model, this will perform the **forward()** method located within the model object

Calculate the **loss value** (how wrong the model's predictions are)

Zero the **optimizer gradients** (they accumulate every epoch, zero them to start fresh each forward pass)

Perform **backpropagation** on the loss function (compute the gradient of every parameter with `requires_grad=True`)

Step the **optimizer** to update the model's parameters with respect to the gradients calculated by `loss.backward()`

Note: all of this can be turned into a function

# PyTorch testing loop

```

1 # Setup empty lists to keep track of model progress
2 epoch_count = []
3 train_loss_values = []
4 test_loss_values = []
5
6 # Pass the data through the model for a number of epochs (e.g. 100) epochs:
7 for epoch in range(epochs):
8     ## Training loop code here ##
9
10    ## Testing starts here ##
11
12    # Put the model in evaluation mode
13    model.eval()
14
15    # Turn on inference mode context manager
16    with torch.inference_mode():
17        # 1. Forward pass on test data
18        test_pred = model(X_test)
19
20        # 2. Calculate loss on test data
21        test_loss = loss_fn(test_pred, y_test)
22
23    # Print out what's happening every 10 epochs
24    if epoch % 10 == 0:
25        epoch_count.append(epoch)
26        train_loss_values.append(loss)
27        test_loss_values.append(test_loss)
28        print(f"Epoch: {epoch} | MAE Train Loss: {loss} | MAE Test Loss: {test_loss}")

```

Create empty lists for storing useful values (helpful for tracking model progress)

Tell the model we want to **evaluate** rather than train (this turns off functionality used for training but not evaluation) (faster performance!)

Turn on `torch.inference_mode()` context manager to disable functionality such as gradient tracking for inference (gradient tracking not needed for inference)

Pass the test data through the model (this will call the model's implemented **forward()** method)

Calculate the **test loss value** (how wrong the model's predictions are on the test dataset, lower is better)

Display **information outputs** for how the model is doing during training/testing every ~10 epochs (note: what gets printed out here can be adjusted for specific problems)

Note: all of this can be turned into a function